

Lab 5

Implement Wordcount program on Hadoop framework

Driver code

```
// Importing libraries import
java.io.IOException; import
org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import
org.apache.hadoop.mapred.FileInputFor
mat; import
org.apache.hadoop.mapred.FileOutputFo
rmat; import
org.apache.hadoop.mapred.JobClient;
import
org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;

public class WCDriver extends Configured implements Tool {

    public int run(String[] args) throws
IOException {        if (args.length < 2) {
        System.out.println("Please give valid
inputs");        return -1;
    }

    JobConf conf = new JobConf(WCDriver.class);
conf.setJobName("WordCount");

    FileInputFormat.setInputPaths(conf, new Path(args[0]));
    FileOutputFormat.setOutputPath(conf, new Path(args[1]));

    conf.setMapperClass(WCMapper.class);
    conf.setReducerClass(WCReducer.class);

    conf.setMapOutputKeyClass(Text.class);
    conf.setMapOutputValueClass(IntWritable.class);

    conf.setOutputKeyClass(Text.class);
    conf.setOutputValueClass(IntWritable.class);

    JobClient.runJob(conf);
    return 0;
}
```

```

// Main Method
public static void main(String[] args) throws
Exception {      int exitCode =
ToolRunner.run(new WCDriver(), args);
    System.out.println("Job Exit Code: " + exitCode);
}
}

```

Mapper Code

```

// Importing libraries import
java.io.IOException; import
org.apache.hadoop.io.IntWritable;
import
org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import
org.apache.hadoop.mapred.MapReduce
Base; import
org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;

public class WCMapper extends MapReduceBase implements Mapper<LongWritable,
Text, Text, IntWritable> {

    // Map function
    public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable>
output, Reporter reporter)      throws IOException {
        String line = value.toString();

        // Splitting the line on whitespace      for
        (String word : line.split("\\s+")) {      if
        (word.length() > 0) {
            output.collect(new Text(word), new
            IntWritable(1));
        }
    }
}
}

```

Reducer Code

```

// Importing libraries
import
java.io.IOException;
import java.util.Iterator;

```

```

import
org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import
org.apache.hadoop.mapred.MapReduce
Base; import
org.apache.hadoop.mapred.OutputColle
ctor; import
org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;

public class WCReducer extends MapReduceBase implements Reducer<Text,
IntWritable, Text, IntWritable> {

    // Reduce function
    public void reduce(Text key, Iterator<IntWritable> values,
        OutputCollector<Text,
IntWritable> output,          Reporter
reporter) throws IOException {      int
count = 0;

        // Counting the frequency of each
word      while (values.hasNext()) {
            count += values.next().get();
        }

        output.collect(key, new IntWritable(count));
    }
}

```

```

HDFS: Number of bytes written=69
HDFS: Number of read operations=15
HDFS: Number of large read operations=0
HDFS: Number of write operations=4
HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Map input records=5
  Map output records=20
  Map output bytes=169
  Map output materialized bytes=215
  Input split bytes=87
  Combine input records=0
  Combine output records=0
  Reduce input groups=10
  Reduce shuffle bytes=215
  Reduce input records=20
  Reduce output records=10
  Spilled Records=40
  Shuffled Maps=1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=0
  Total committed heap usage (bytes)=1052770304
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=89
File Output Format Counters
  Bytes Written=69
0
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC: $ hadoop fs -ls /output/
ls: /output/: No such file or directory
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC: $ hadoop fs -ls /output
ls: /output: No such file or directory
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC: $ hadoop fs -ls output
Found 2 items
-rw-r--r-- 1 hadoop supergroup          0 2025-04-29 15:34 output/_SUCCESS
-rw-r--r-- 1 hadoop supergroup        69 2025-04-29 15:34 output/part-00000
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC: $ hadoop fs -cat output/part-00000
are 1
brother 1
family 1
hi 1
how 5
ls 4
job 1
sister 1
you 1
your 4
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC: $

```